


Priority-Based Cache Allocation in Throughput Processors

Dong Li*, Minsoo Rhu*, Daniel R. Johnson, Mike O'Connor, Mattan Erez,
Doug Burger, Donald S. Fussell and Stephen W. Keckler



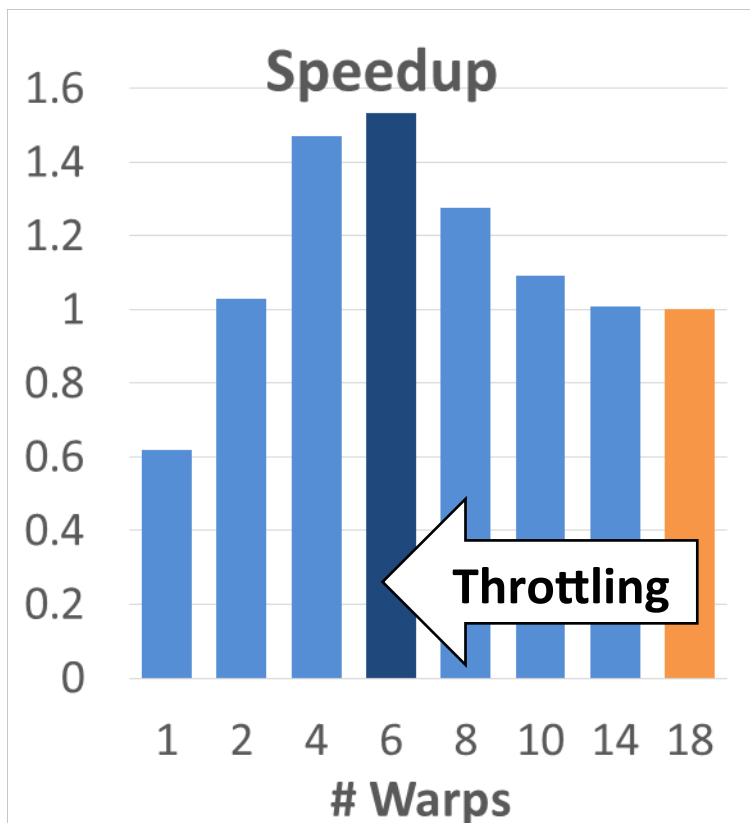
*First authors Li and Rhu have made equal contributions to this work and are listed alphabetically

Introduction


- Graphics Processing Units (GPUs)
 - General purpose throughput processors
 - Massive thread hierarchy, warp=32 threads
- **Challenge:**
 - Small caches + many threads → contention + cache thrashing
- Prior work: throttling thread level parallelism (TLP)
 - **Problem 1:** under-utilized system resources
 - **Problem 2:** unexplored cache efficiency

Motivation: Case Study (CoMD)


Throttling improves performance



Observation 1: Resource under-utilized




Observation 2: Unexplored cache efficiency



■ Best Throttled ■ max TLP (default)

Motivation


- Throttling
 - Tradeoff between cache efficiency and parallelism.
 - $\# \text{ Active Warps} = \# \text{ Warps Allocating Cache}$
- Idea
 - Decoupling cache efficiency from parallelism



Our Proposal: Priority Based Cache Allocation (PCAL)

Baseline


Standard operation



- ▶ Regular threads: threads have all capabilities, operate as usual (full access to the L1 cache)

TLP Throttling

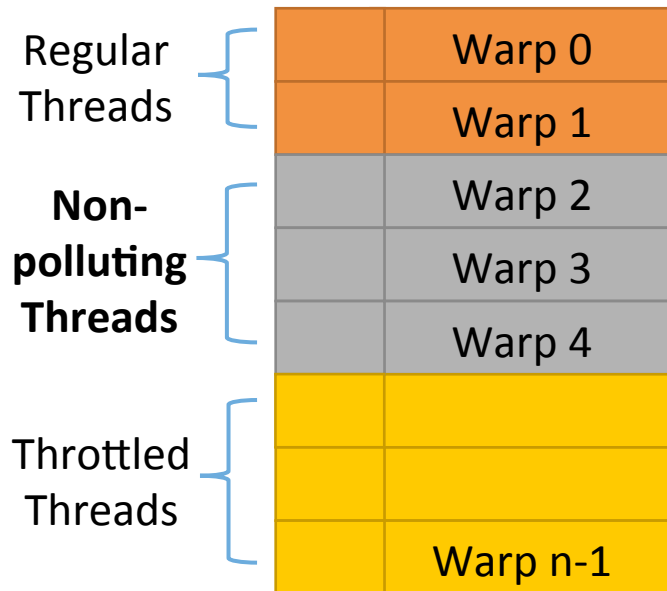
TLP reduced to improve performance



- ▶ Regular threads: threads have all capabilities, operate as usual (full access to the L1 cache)
- ▶ Throttled threads: throttled (not runnable)

Proposed Solution: PCAL

Three subsets of threads



- ▶ Regular threads: threads have all capabilities, operate as usual (full access to the L1 cache)
- ▶ Non-polluting threads: runnable, but prevented from *polluting* L1 cache
- ▶ Throttled threads: throttled (not runnable)

Proposed Solution: PCAL

Tokens for privileged cache access

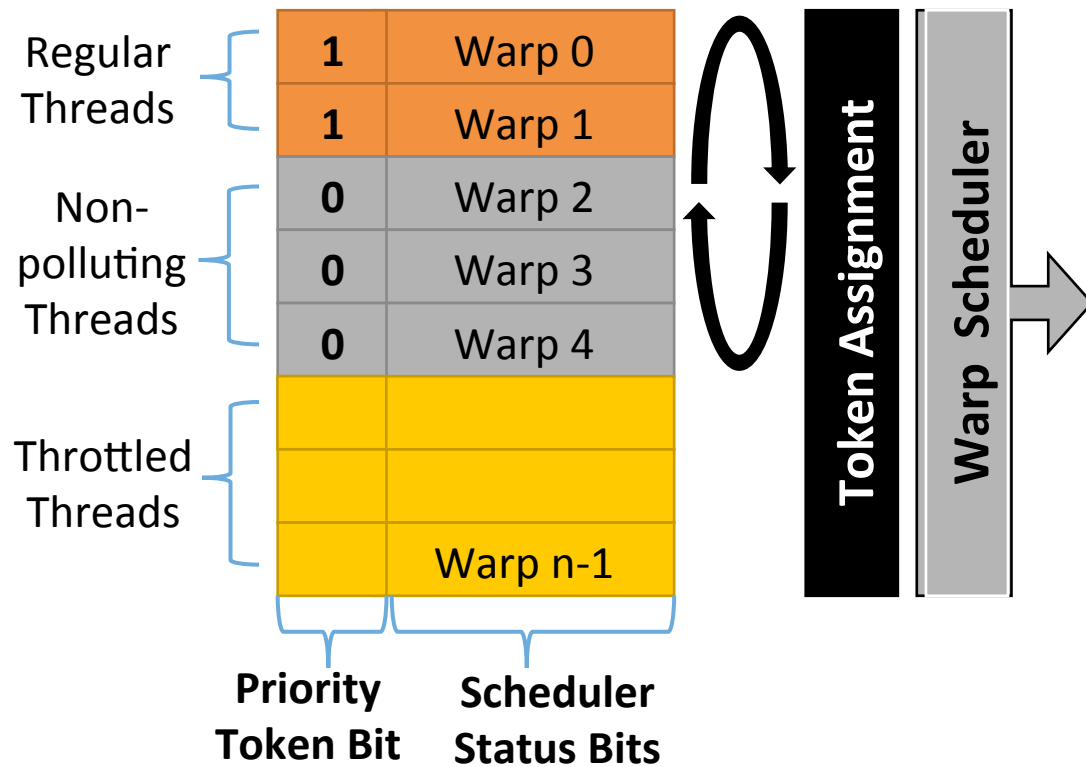
Regular Threads	1	Warp 0
	1	Warp 1
Non-polluting Threads	0	Warp 2
	0	Warp 3
	0	Warp 4
Throttled Threads		
		Warp n-1

Priority Token Bit

- ▶ **Tokens** grant capabilities or privileges to some threads
- ▶ Threads with tokens are allowed to **allocate and evict** L1 cache lines
- ▶ Threads without tokens can read and write, but **not allocate or evict**, L1 cache lines

Proposed solution: Static PCAL

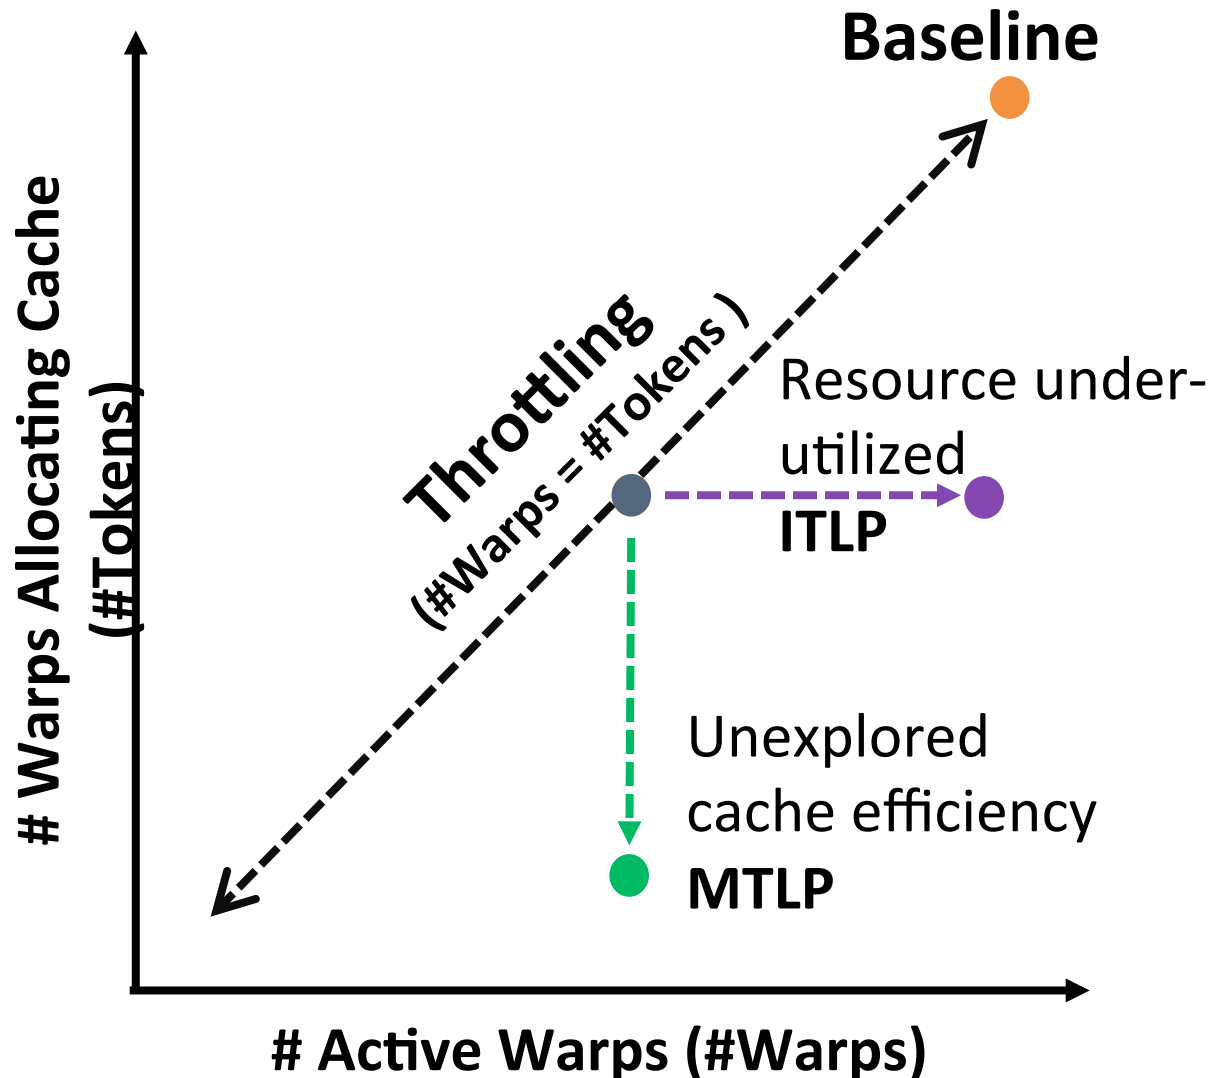
Enhanced scheduler operation



- ▶ HW implementation: simplicity
 - ▶ Scheduler manages per-warp token bits
 - ▶ Token bits sent with memory requests
- ▶ Policies: token assignment
 - ▶ Assign at warp launch, release at termination
 - ▶ Re-assigned to oldest token-less warp
- ▶ Parameters supplied by software prior to kernel launch

Two Optimization Strategies

to exploit the 2-D (#Warps, #Tokens) search space



- **1. Increasing TLP (ITLP)**
 - Adding non-polluting warps
 - Without hurting regular warps
- **2. Maintaining TLP (MTLP)**
 - Reduce #Token to increase the efficiency
 - Without decreasing TLP

Proposed Solution: Dynamic PCAL

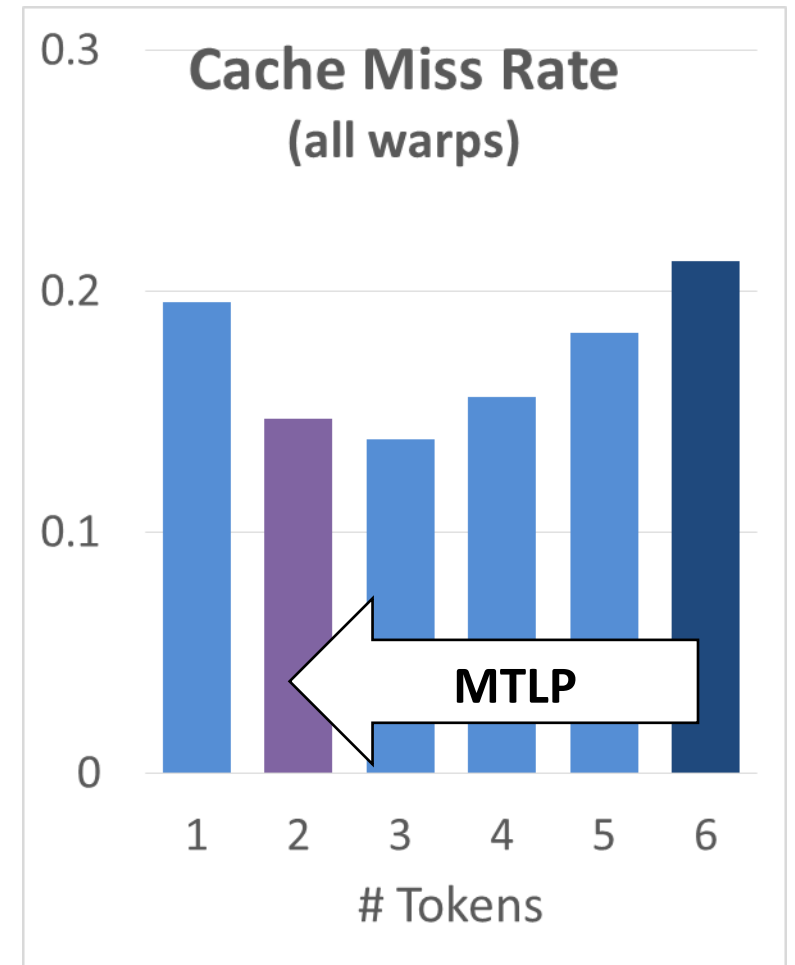


- Decide parameters at run time
 - Choose MTLP/ITLP based on resource usage performance counter
 - For MTLP: search **#Tokens** in parallel
 - For ITLP: search **#Warps** in sequential
- Please refer to our paper for more details

Evaluation

- ▶ Simulation: GPGPU-Sim
 - ▶ Open source academic simulator, configured similar to Fermi (full-chip)
- ▶ Benchmarks
 - ▶ Open source GPGPU suites: Parboil, Rodinia, LonestarGPU. etc
 - ▶ Selected benchmark subset shows *sensitivity to cache size*

MTLP Example (CoMD)

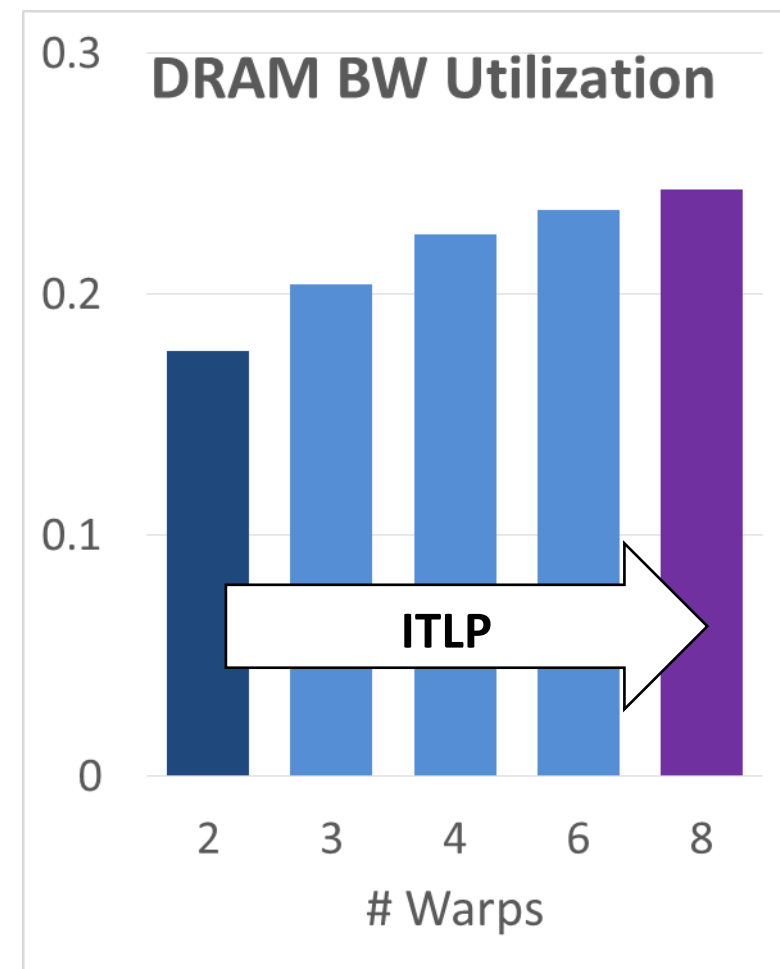

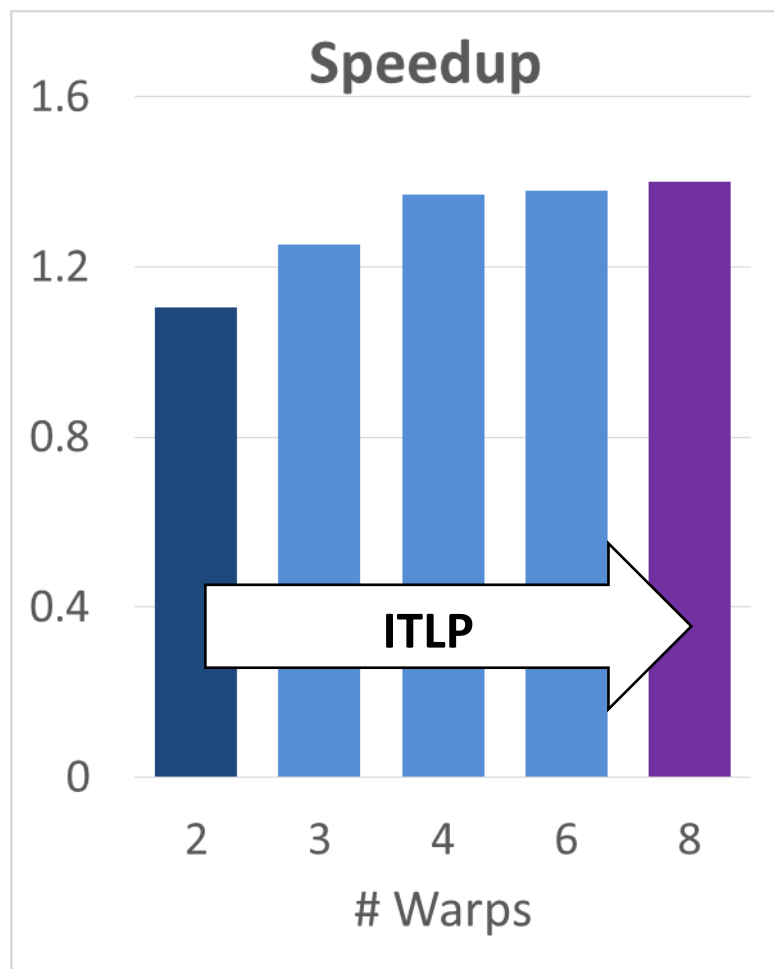
MTLP: reducing #Tokens, keeping #Warps =6



■ PCAL with MTLP ■ Best Throttled

ITLP Example (Similarity-Score)


ITLP: increasing #Warps, keeping #Tokens=2



■ PCAL with ITLP

■ Best Throttled

PCAL Results Summary



- ▶ Baseline: best throttled results
- ▶ Performance improvement: Static-PCAL: 17%, Dynamic-PCAL: 11%

Conclusion

- ▶ Existing throttling approach may lead to two problems:
 - ▶ Resources underutilization
 - ▶ Unexplored cache efficiency
- ▶ We propose PCAL, a simple mechanism to rebalance cache efficiency and parallelism
 - ▶ MTLP: alleviates cache thrashing while maintaining parallelism
 - ▶ ITLP: increases parallelism without hurting regular warps
- ▶ Throughput improvement over best throttled results
 - ▶ Static PCAL: 17%
 - ▶ Dynamic PCAL: 11%