



# Maximizing SIMD Resource Utilization in GPGPUs with SIMD Lane Permutation



**Minsoo Rhu and Mattan Erez**

**The University of Texas at Austin  
Electrical and Computer Engineering Department**



# Introduction

- SIMD-based GPU Architectures
  - Suffers from low SIMD efficiency when application contains highly irregular control flow
  - **Compaction**-based architectures have recently gained high interest as a way of minimizing the waste in SIMD units
    - *Problem*: applicability of compaction limited to highly divergent applications, despite its design overhead
- Main contribution of this paper
  - Analysis on the cause of limited effectiveness of compaction
    - Concentration of active threads to particular SIMD lanes
  - Static/deterministic permutation of thread assignment to SIMD units to enhance the effectiveness of compaction
    - Improve SIMD utilization and performance
    - Widen the applicability of compaction schemes




# Outline

- **GPU and SIMD compaction background**
- Aligned divergence and compactability
- SLP – SIMD Lane Permutation
- Evaluation

# Graphic Processing Units (GPUs)

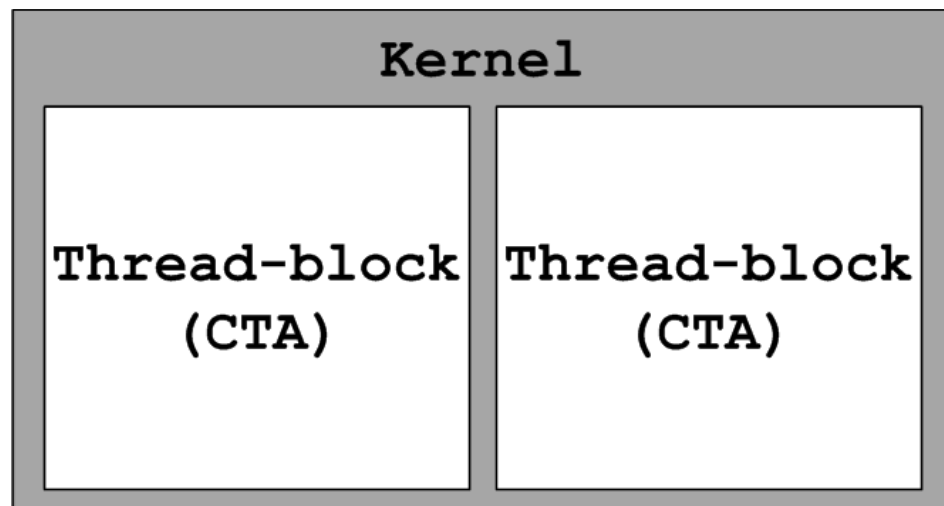
- General-purpose many-core accelerators
  - Supports non-graphics APIs (e.g. CUDA, OpenCL)
- Scalar frontend (fetch & decode) + parallel backend
  - Amortizes the cost of frontend and control





# CUDA exposes hierarchy of data-parallel threads

- **SPMD model**: single **kernel** executed by all threads
- **Kernel / Thread-block**
  - Multiple **thread-blocks** (concurrent-thread-arrays(**CTAs**)) compose a **kernel**






# CUDA exposes hierarchy of data-parallel threads

- **SPMD model**: single **kernel** executed by all threads
- **Kernel / Thread-block / Warp / Thread**
  - Multiple **warps** compose a **thread-block**
  - Multiple **threads (32)** compose a warp

**A warp is scheduled as a *batch* of threads**



**Assumption: width of SIMD lanes matches warp size (i.e. 32 SIMD lanes)**

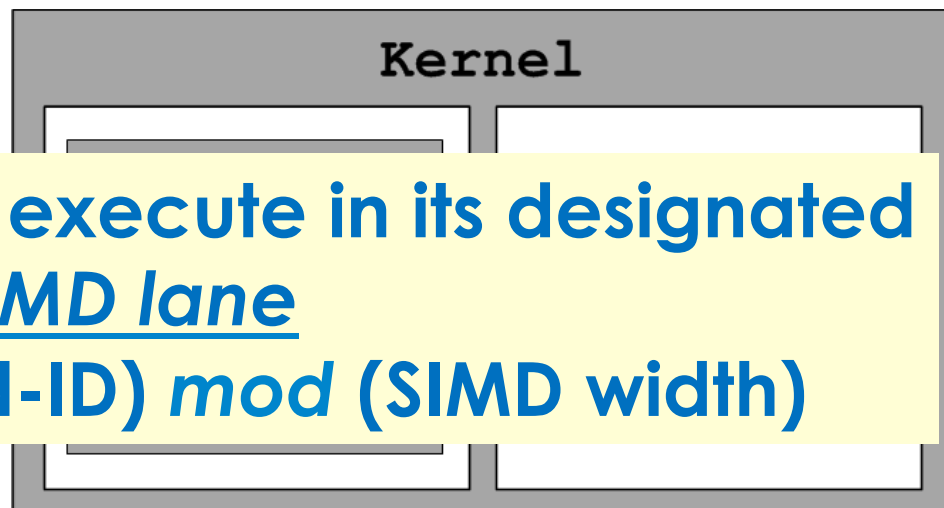



# CUDA exposes hierarchy of data-parallel threads

- **SPMD model**: single **kernel** executed by all threads
- **Kernel / Thread-block / Warp / Thread**
  - Multiple **warps** compose a **thread-block**
  - Multiple **threads (32)** compose a warp

: Thread-ID 0, 32, 64 ... execute in physical lane #0

: Thread-ID 2, 34, 66 ... execute in physical lane #2




Threads execute in its designated  
home SIMD lane  
: (thread-ID) mod (SIMD width)

Assumption: width of SIMD lanes matches warp size (i.e. 32 SIMD lanes)



# GPUs have HW support for conditional branches

- **Control divergence:** threads in warp branch differently
  - a.k.a. **Branch Divergence**
  - *Stack-based divergence/reconvergence model*
  - Active bitmasks (or predicate bitmasks):
    - Bitmasks for both true/false paths dynamically derived
    - Allows subset of a warp to commit results




(a) Example Control Flow Graph

1: Active, 0: Inactive






# Challenge: Underutilization of SIMD units



(a) Example Control Flow Graph



(b) Execution flow

Time



: Threads (lanes) masked out from execution, remaining idle.

- Number of active threads in a warp decreases every time control diverges
- Active area of research
  - Thread block compaction [Fung'11]
  - Larger warps [Narasiman'11]




# Thread-block compaction (TBC) [Fung'11]


- Dynamically **compact** warps within a **thread-block**
  - Synchronize all warps at conditional-branches and reconvergence points
  - Threads with different home SIMD lanes compacted together

: [0,1... A,B] refers to active thread-IDs executing in that basic block

: [-] refers to inactive threads, masked out from execution



(a) Example control flow graph



(b) Execution flow without compaction




# Thread-block compaction (TBC) [Fung'11]


- Dynamically **compact** warps within a **thread-block**
  - Synchronize all warps at conditional-branches and reconvergence points
  - Threads with different home SIMD lanes compacted together

: [0,1... A,B] refers to active thread-IDs executing in that basic block

: [-] refers to inactive threads, masked out from execution



(a) Example control flow graph



(b) Execution flow without compaction




# Thread-block compaction (TBC) [Fung'11]


- Dynamically **compact** warps within a **thread-block**
  - Synchronize all warps at conditional-branches and reconvergence points
  - Threads with different home SIMD lanes compacted together

: [0,1... A,B] refers to active thread-IDs executing in that basic block


: [-] refers to inactive threads, masked out from execution



(a) Example control flow graph



(b) Execution flow without compaction



(c) Execution flow with compaction




# Thread-block compaction (TBC) [Fung'11]


- Dynamically **compact** warps within a **thread-block**
  - Synchronize all warps at conditional-branches and reconvergence points
  - Threads with different home SIMD lanes compacted together

: [0,1... A,B] refers to active thread-IDs executing in that basic block


: [-] refers to inactive threads, masked out from execution



(a) Example control flow graph



(b) Execution flow without compaction




(c) Execution flow with compaction



## TBC as-is

- Average SIMD lanes occupied for execution
  - No\_TBC** : Baseline architecture *without* compaction
  - TBC** : baseline compaction mechanism

Higher is better






## TBC as-is

- Average SIMD lanes occupied for execution
  - No\_TBC** : Baseline architecture without compaction

**TBC, in general, only effective for highly divergent applications**

Higher is better





# Outline

- GPU and SIMD compaction background
- **Branch compactability and aligned divergence**
- SLP – SIMD Lane Permutation
- Evaluation





# When Does Compaction Fail?

- Most (or all) of SIMD lanes already occupied
  - Compaction inherently impossible

Active Threads	
$W_0$	0123
$W_1$	----
$W_2$	89AB
$W_3$	----

Active Threads	
$W_0$	012-
$W_1$	456-
$W_2$	89A-
$W_3$	CDE-

Active Threads	
$W_0$	0123
$W_1$	4567
$W_2$	89AB
$W_3$	CD--



# When Does Compaction Fail?

- Most (or all) of SIMD lanes already occupied
  - Compaction inherently impossible

Active Threads	
$W_0$	0123
$W_1$	----
$W_2$	89AB
$W_3$	----

Active Threads	
$W_0$	012-
$W_1$	456-
$W_2$	89A-
$W_3$	CDE-

Active Threads	
$W_0$	0123
$W_1$	4567
$W_2$	89AB
$W_3$	CD--



# When Does Compaction Fail?

- Most (or all) of SIMD lanes already occupied
  - Compaction inherently impossible

Active Threads	
$W_0$	0123
$W_1$	----
$W_2$	89AB
$W_3$	----

Active Threads	
$W_0$	012-
$W_1$	456-
$W_2$	89A-
$W_3$	CDE-

Active Threads	
$W_0$	0123
$W_1$	4567
$W_2$	89AB
$W_3$	CD--

- Active threads aligned (clustered) on certain lanes
  - Compaction theoretically feasible, if crossbars can distribute operands across different SIMD lanes

Active Threads	
$W_0$	0----
$W_1$	4----
$W_2$	8----
$W_3$	C----

Active Threads	
$W_0$	01--
$W_1$	45--
$W_2$	89--
$W_3$	CD--

Active Threads	
$W_0$	0-2-
$W_1$	4-6-
$W_2$	8-A-
$W_3$	C-E-



# When Does Compaction Fail?

- Most (or all) of SIMD lanes already occupied
  - Compaction inherently impossible

Active Threads	
W <sub>0</sub>	0123
W <sub>1</sub>	----
W <sub>2</sub>	89AB
W <sub>3</sub>	----

Active Threads	
W <sub>0</sub>	012-
W <sub>1</sub>	456-
W <sub>2</sub>	89A-
W <sub>3</sub>	CDE-

Active Threads	
W <sub>0</sub>	0123
W <sub>1</sub>	4567
W <sub>2</sub>	89AB
W <sub>3</sub>	CD--

- Active threads aligned (clustered) on certain lanes
  - Compaction then fails due to **Aligned Divergence!** (operands across different SIMD lanes)

Active Threads	
W <sub>0</sub>	0----
W <sub>1</sub>	4----
W <sub>2</sub>	8----
W <sub>3</sub>	C----

Active Threads	
W <sub>0</sub>	01--
W <sub>1</sub>	45--
W <sub>2</sub>	89--
W <sub>3</sub>	CD--

Active Threads	
W <sub>0</sub>	0-2-
W <sub>1</sub>	4-6-
W <sub>2</sub>	8-A-
W <sub>3</sub>	C-E-



## Aligned divergence – (1)

- Case 1: Branch condition depends on **data** array
  - Each thread references **different** element of the array
  - Unlikely to cause aligned divergence

---

**Code #1)** Branch depending on data arrays

---

```
0    // Code snippet from the kernel of BFS benchmark
1    // g_graph_visited and g_graph_edges are data array parameters.
2
3    int tid = blockIdx.x*MAX_THREADS_PER_BLOCK + threadIdx.x;
4
5    ...
6    int id = g_graph_edges[...];
7    if( !g_graph_visited[id] )
8    {
9        ...
10 }
```

---



## Aligned divergence – (1)

- Case 1: Branch condition depends on **data** array
  - Each thread references **different** element of the array
  - Unlikely to cause aligned divergence

---

**Code #1)** Branch depending on data arrays

---

```
0 // Code snippet from the kernel of BFS benchmark
1 // g_graph_visited and g_graph_edges are data array parameters.
2
3 int tid = blockIdx.x*MAX_THREADS_PER_BLOCK + threadIdx.x;
4
5 ...
6 int id = g_graph_edges[...];
7 if( !g_graph_visited[id] )
8 {
9     ...
10 }
```

**D-Branches**



## Aligned divergence – (2)

- Case 2: Branch cond. depends on **programmatic** value
  - *Indices of thread-ID, warp-ID, CTA-ID, width/height of CTA*
  - *Scalar input parameters to the kernel, constants*
  - Threads sharing home SIMD lane likely to reference **same** value

---

**Code #2)** Programmatic branch causing only the 1<sup>st</sup> half of the warp active

---

```
0    // Code snippet from the kernel of BACKP benchmark
1    // CTA is a (8 × 16) 2-D array of threads.
2
3    int tx = threadIdx.x;
4    int ty = threadIdx.y;
5    ...
6    for (int i=1; i<=__log2f(HEIGHT); i++){
7        int power_two = __powf(2,i);
8
9        if( ty % power_two == 0 ) { ... }
10       ...
11    }
```



## Aligned divergence – (2)

- Case 2: Branch cond. depends on **programmatic** value
  - *Indices of thread-ID, warp-ID, CTA-ID, width/height of CTA*
  - *Scalar input parameters to the kernel, constants*
  - Threads sharing home SIMD lane likely to reference **same** value

---

**Code #2)** Programmatic branch causing only the 1<sup>st</sup> half of the warp active

---

```
0    // Code snippet from the kernel of BACKP benchmark
1    // CTA is a (8 × 16) 2-D array of threads.
2
3    int tx = threadIdx.x;
4    int ty = threadIdx.y;
5    ...
6    for (int i=1; i<=__log2f(HEIGHT); i++){
7        int power_two = __powf(2,i);
8
9        if( ty % power_two == 0 ) { ... }
10       ...
11    }
```

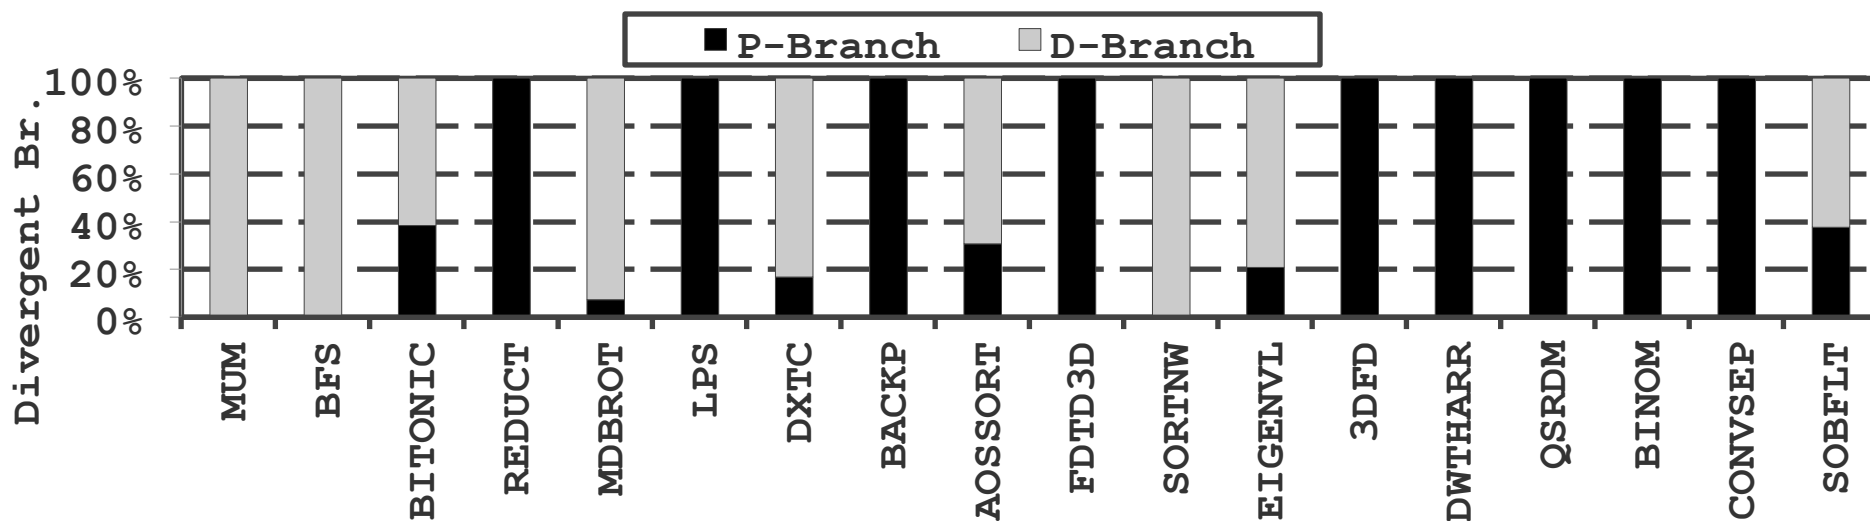
**P-Branches**





# Compaction Rate of P-/D-Branches (with TBC)

- Definition:** Fraction of compactable paths among *all* paths generated by divergent branches




Branch categorization done with GPUOcelot  
using Taint-analysis (detailed in paper)



# Compaction Rate of P-/D-Branches (with TBC)

- Definition:** Fraction of compactable paths among *all* paths generated by divergent branches




Ideal: TBC with crossbar

Higher is better



## Compaction Rate of P / D Branches (with TBC)

**D-branch compacts well with TBC**  
**P-branch not so much compared to *Ideal***



**Ideal: TBC with crossbar**


**Higher is better**



# TBC with Crossbar (*Ideal* Compaction)

- Average SIMD lanes occupied for execution

Higher is better







Significant opportunities to  
enhance SIMD efficiency with P-branches.

But crossbars are expensive!

tion)  
ecution

Higher is better





# Outline


- GPU and SIMD compaction background
- Branch compactability and aligned divergence
- **SLP – SIMD Lane Permutation**
- Evaluation



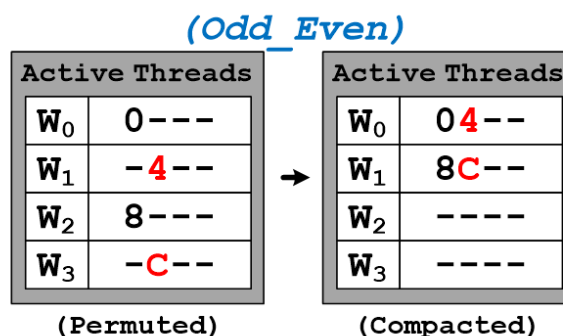
# SIMD Lane Permutation (SLP)

- Motivation:** Permute home SIMD lanes from their sequentially assigned location to alleviate aligned divergence

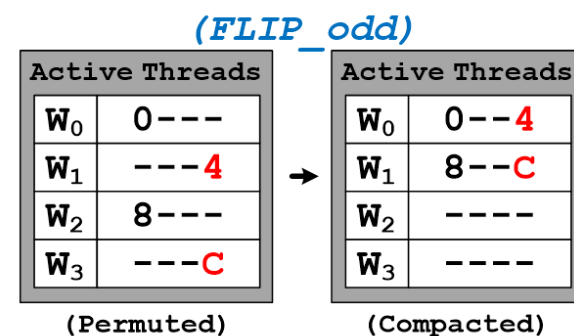
\* WID: Warp-ID



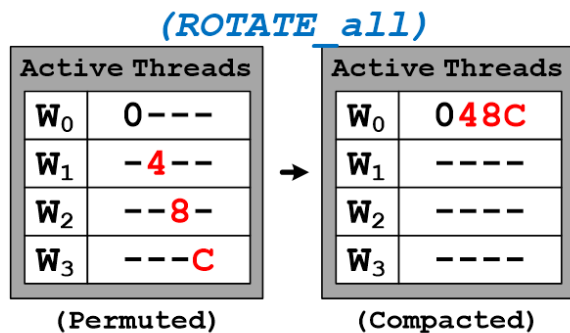
(a) Compaction *as-is*,



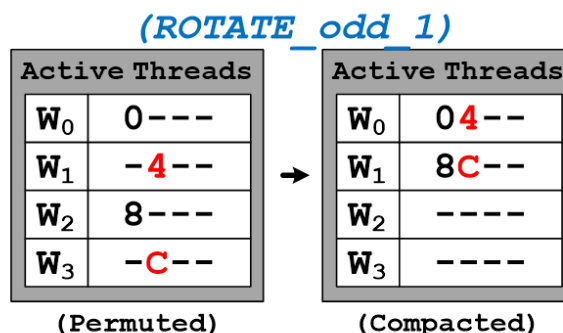
(b) XOR warps with  
odd WID by 1



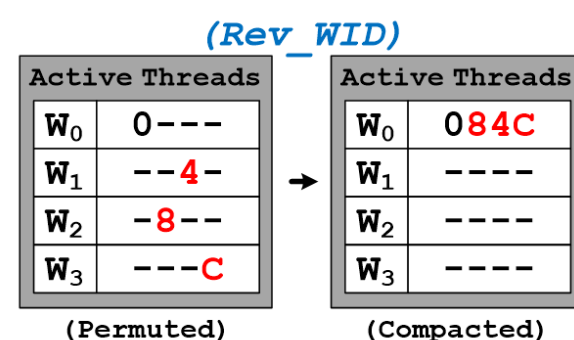
(c) Flip warps with  
odd WID



(d) Rotate *all* warps  
by WID



(e) Rotate warps with  
odd WID by 1



(f) XOR all warps by  
bit-reverse of WID

- Odd\_Even: proposed by Fung et al. [MICRO'07] for tuning/optimizing BITONIC sorting application



# Pitfalls of *Randomly* Permuting Lanes

- **Observation:** Many P-branches invoke up to half of lanes active after divergence
- Compaction fails unless active threads permuted exactly to the other half of (vacant) lanes

Active Threads	
$W_0$	0123----
$W_1$	89AB----
$W_2$	GHIJ----
$W_3$	OPQR----

Active Threads	
$W_0$	01--45--
$W_1$	89--CD--
$W_2$	GH--KL--
$W_3$	OP--ST--

Active Threads	
$W_0$	0-2-4-6-
$W_1$	8-A-C-E-
$W_2$	G-I-K-M-
$W_3$	O-Q-S-U-





# Balanced Permutation

- **Observation:** P-branches frequently exhibit highly skewed, predictable concentration of active lanes
- **Intuition** : Distribute active threads across SIMD lanes evenly in a balanced manner

Lane-ID	0	1	2	3	4	5	6	7
	000	001	010	011	100	101	110	111
For W0	0	1	2	3	4	5	6	7
For W1	0	1	2	3	4	5	6	7
For W2	0	1	2	3	4	5	6	7
For W3	0	1	2	3	4	5	6	7
For W4	0	1	2	3	4	5	6	7
For W5	0	1	2	3	4	5	6	7
For W6	0	1	2	3	4	5	6	7
For W7	0	1	2	3	4	5	6	7



## Balanced Permutation

- **Observation:** P-branches frequently exhibit highly skewed, predictable concentration of active lanes
- **Intuition** : Distribute active threads across SIMD lanes evenly in a balanced manner

Lane-ID	0	1	2	3	4	5	6	7
	000	001	010	011	100	101	110	111
For W0	0	1	2	3	4	5	6	7
For W1	0	1	2	3	4	5	6	7
For W2	0	1	2	3	4	5	6	7
For W3	0	1	2	3	4	5	6	7
For W4	0	<b>Baseline: assigned lane-IDs based on <u>sequential</u> assignment</b>						
For W5	0							
For W6	0	1	2	3	4	5	6	7
For W7	0	1	2	3	4	5	6	7



# Balanced Permutation

- **Observation:** P-branches frequently exhibit highly skewed, predictable concentration of active lanes
- **Intuition** : Distribute active threads across SIMD lanes evenly in a balanced manner

Lane-ID	0	1	2	3	4	5	6	7
	000	001	010	011	100	101	110	111
For W0	0	1	2	3	4	5	6	7
For W1	0	1	2	3	4	5	6	7
For W2	0	1	2	3	4	5	6	7
For W3	0	1	2	3	4	5	6	7
For W4	0	1	2	3	4	5	6	7
For W5	0	1	2	3	4	5	6	7
For W6	0	1	2	3	4	5	6	7
For W7	0	1	2	3	4	5	6	7



## Balanced Permutation

- Observation:** P-branches frequently exhibit highly skewed.  
**< Balanced permutation algorithm >**
- Intuiti- Even-ID warps:** **XOR** lane-ID by **(Warp-ID >> 1)**  
**- Odd-ID warps:** **XOR** lane-ID by **~(Warp-ID >> 1)**

Lane-ID	0	1	2	3	4	5	6	7
	000	001	010	011	100	101	110	111
For W0	0	1	2	3	4	5	6	7
For W1	0	1	2	3	4	5	6	7
For W2	0	1	2	3	4	5	6	7
For W3	0	1	2	3	4	5	6	7
For W4	0	1	2	3	4	5	6	7
For W5	0	1	2	3	4	5	6	7
For W6	0	1	2	3	4	5	6	7
For W7	0	1	2	3	4	5	6	7



# Balanced Permutation

- Observation:** P-branches frequently exhibit highly skewed.  
**< Balanced permutation algorithm >**
- Intuiti-** Even-ID warps: **XOR lane-ID by (Warp-ID >> 1)**  
- Odd-ID warps: **XOR lane-ID by  $\sim$ (Warp-ID >> 1)**

	Lane-ID	0	1	2	3	4	5	6	7
		000	001	010	011	100	101	110	111
For W0	XOR-000	0	1	2	3	4	5	6	7
For W1	XOR-111	7	6	5	4	3	2	1	0
For W2	XOR-001	1	0	3	2	5	4	7	6
For W3	XOR-110	6	7	4	5	2	3	0	1
For W4	XOR-010	2	3	0	1	6	7	4	5
For W5	XOR-101	5	4	7	6	1	0	3	2
For W6	XOR-011	3	2	1	0	7	6	5	4
For W7	XOR-100	4	5	6	7	0	1	2	3



# Outline

- GPU and SIMD compaction background
- Branch compactability and aligned divergence
- SLP – SIMD Lane Permutation
- **Evaluation**



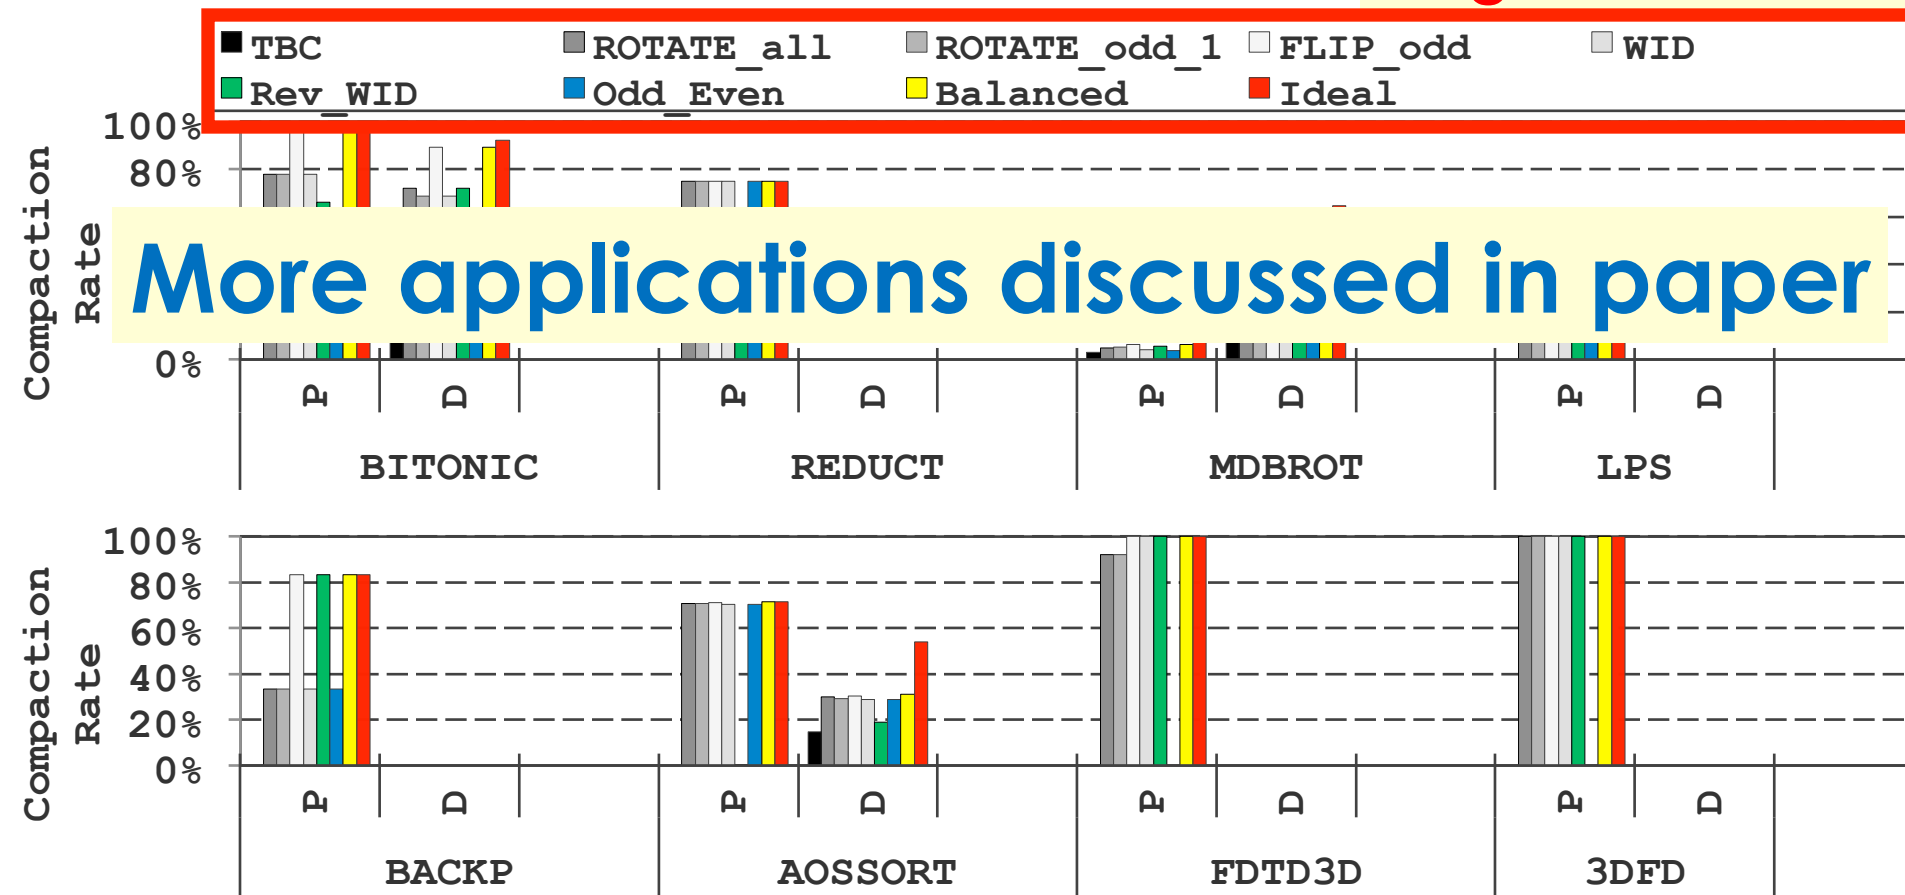
# Simulation Environment

- GPUOcelot (r1865) – based on TBC
  - Branch categorization
  - Compaction rate
  - SIMD lane utilization
- GPGPU-Sim (v2.1) – based on TBC + CAPRI\*
  - Performance study
  - Similar to QuadroFX-5800
  - uArch configurations detailed in paper
- Workloads
  - Chosen from CUDA-SDK(v3.0), Rodinia, Parboil, etc

# Compaction Rate

- **Definition:** Fraction of compactable paths among all paths generated by divergent branches
- **SLP:** significantly improves *P*-branch compaction rate

# Higher is better












# Average SIMD Lane Utilization

- Definition:** average number of SIMD lanes *actually* executing and committing results
- Compaction rate doesn't tell '*how effectively*' warps are reduced


Higher is better





# Average SIMD Lane Utilization


- **Definition:** average number of SIMD lanes actually executing
- **TBC:** lowest utilization due to un-compacted P-branches
- **SLP:**
  - **Odd\_Even:** average 2.1% (max 18%) increase over TBC
  - **Balanced:** average 7.1% (max 34%) increase over TBC





# Speedup

- **Baseline:** No compaction
- TBC+CAPRI only effective for 'irregular' applications
- SLP widens the range of applications that benefit from compaction techniques
  - **7.1%** (max 34%) improvements on top of TBC+CAPRI





# Conclusions

- Effectiveness of previous compactions limited to highly irregular applications
  - Only effective for D-branches
  - Non-compactable P-branches
- SLP enables instrumenting such lost opportunities
  - Requires simple re-mapping of thread-IDs to Lane-ID
  - Enhance compaction rate and SIMD utilization
- Throughput improvements
  - **7.1%** (max 34%) improvements on top of TBC+CAPRI